

---

# Reinforcement Learning

---

**Zhuoyuan Chen\***  
Facebook AI Research  
Menlo Park, CA 94025  
chenzhuoyuan07@gmail.com

## Abstract

RL: Policy Gradient, Q-Learning, actor-critic

## 1 Policy Gradient

Reward, trajectories:  $\tau$ :

$$J(\theta) = \mathbb{E}_{\tau \sim p_{\theta}(\tau)} \left[ \sum_t r(s_t, a_t) \right] \approx \frac{1}{N} \sum_i \sum_t r(s_{i,t}, a_{i,t}) \quad (1)$$

where

$$J(\theta) = \int \pi_{\theta}(\tau) r(\tau) d\tau$$

Then

$$\nabla_{\theta} J(\theta) = \int \pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau) d\tau = \mathbb{E}_{\tau} [\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]$$

1. Causal in  $r$ ;
2. Minus baseline to reduce variance, minimum variance achieved at

$$b = \frac{\mathbb{E}[g(\tau)^2 r(\tau)]}{\mathbb{E}[g(\tau)^2]}$$

3. On policy;

Importance sampling:

$$\mathbb{E}_{x \sim p(x)} [f(x)] = \mathbb{E}_{x \sim q(x)} \left[ \frac{p(x)}{q(x)} f(x) \right]$$

Given a trajectory, we have

$$\frac{\pi_{\theta}(\tau)}{\bar{\pi}_{\theta}(\tau)} = \frac{\prod_{t=1}^T \pi_{\theta}(a_t | s_t)}{\prod_{t=1}^T \bar{\pi}_{\theta}(a_t | s_t)}$$

Then off-policy PG with IS:

$$\nabla_{\theta'} J(\theta') = \mathbb{E}_{\tau \sim \pi_{\theta}(\tau)} \left[ \sum_{t=1}^T \nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \left( \prod_{t'=1}^t \frac{\pi_{\theta'}(a_{t'} | s_{t'})}{\pi_{\theta}(a_{t'} | s_{t'})} \right) \left( \sum_{t''=t}^T r(s_{t'}, a_{t'}) \prod_{t''=t}^{t'} \frac{\pi_{\theta'}(a_{t''} | s_{t''})}{\pi_{\theta}(a_{t''} | s_{t''})} \right) \right] \quad (2)$$

---

\*

## 1.1 TRPO

$$\begin{aligned}
J(\theta') - J(\theta) &= J(\theta') - \mathbb{E}_{s_0 \sim p(s)}[V^{\pi_{\theta}}(s_0)] \\
&= J(\theta') - \mathbb{E}_{\tau \sim p_{\theta'}(\tau)}[V^{\pi_{\theta}}(s_0)] \\
&= \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \sum_{t=1}^{\infty} \gamma^t r(s_t, a_t) \right] + \left[ \sum_{t=1}^{\infty} \gamma^t (\gamma V^{\pi_{\theta}}(s_{t+1}) - V^{\pi_{\theta}}(s_t)) \right] \\
&= \mathbb{E}_{\tau \sim p_{\theta'}(\tau)} \left[ \sum_{t=1}^{\infty} \gamma^t A^{\pi_{\theta}}(s_t, a_t) \right]
\end{aligned}$$

## 1.2 PPO

Loss: let

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{old}}(a_t | s_t)}$$

Then,

$$L^{CLIP}(\theta) = \mathbb{E}[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t)] \quad (3)$$

**PPO system design:** same as A2C (check section 2.1 actor-critic), only difference:

```

step 4: agent.update()
    ratio = exp(action_log_prob - old_action_log_prob)
    surr1 = ratio * adv_target
    surr2 = clamp(ratio, 1.0-eps, 1.0+eps) * adv_target
    action_loss = -min(surr1, surr2)

```

## 1.3 ACKTR

1. Paper: Scalable trust-region method for deep reinforcement learning using Kronecker-factored approximation [NIPS 2017];
2. Use the **K-FAC** technique to approximate Fisher Information matrix in paper [J. Martens and R. Grosse. Optimizing neural networks with kronecker-factored approximate curvature.]

Important properties of Kronecker operator  $\otimes$ :

1.  $(P \otimes Q)^{-1} = P^{-1} \otimes Q^{-1}$ ;
2.  $(P \otimes Q)\text{vec}(T) = PTQ^T$

Let  $W \in \mathbb{R}^{C_{out} \times C_{in}}$  be the weight matrix of the  $l^{th}$  layer, with  $s = Wa$ , we have K-FAC:

$$F_l = \mathbb{E}[\text{vec}\{\nabla_W L\} \text{vec}\{\nabla_W L\}^T] = \mathbb{E}[aa^T \otimes \nabla_s L(\nabla_s L)^T] \quad (4)$$

$$\approx \mathbb{E}[aa^T] \otimes \mathbb{E}[\nabla_s L(\nabla_s L)^T] = A \otimes S := \hat{F}_l \quad (5)$$

Then,

$$\text{vec}(\Delta W) = \hat{F}_l^{-1} \text{vec}\{\nabla_W J\} = \text{vec}(A^{-1} \nabla_W J S^{-1}) \quad (6)$$

In actor-critic setting (two output heads— action  $a$  and critic  $v$ ), we will try to apply  $\mathbb{E}_{p(\tau)}[\nabla \log p(a, v|s) \nabla \log p(a, v|s)^T]$ .

Step-size, following TRPO, SGD-like  $\theta \leftarrow \theta - \eta F^{-1} \nabla_{\theta} L$  will result in large updates in policy and make algorithm prematurely converge to near-deterministic. Therefore, we adopt step-size as:

$$\eta = \min(\eta_{\max}, \sqrt{\frac{2\delta}{\Delta \theta^T \hat{F} \Delta \theta}}) \quad (7)$$

System design:

1. Agent (ACKTR) initialization:

- (a) Split modules (bias)
- (b) Register forward pre-hook (save input)
- (c) Register backward hook (save gradient output)
- (d) Running stats:

```

self.m_aa = {} # save input.t() @ input
self.m_gg = {}
self.Q_a, self.Q_g = {}, {}
self.d_a, self.d_g = {}, {}

```

2. Agent (ACKTR) update:

- (a) Before forward: compute running stat of  $a^T a$  and save in m-aa;
- (b) During backward: compute running stat of  $g^T g$

3. Optimizer.step() for each module:

- (a) Eig SVD of inputs:  $Cov(a, a) = Q_a D_a Q_a^T$ ;
- (b) Eig SVD of output gradient:  $Cov(g, g) = Q_g D_g Q_g^T$ ;
- (c) Update with normal gradient G and update each module's gradient with  $\Delta\theta$ :

$$\Delta\theta = Q_g \frac{Q_g^T G Q_a}{d_g^T d_a + 0.01} Q_a^T \quad (8)$$

- (d) Normal SGD;

## 2 Actor-Critic

value function, advantage:

$$\begin{aligned}
Q^\pi(s_t, a_t) &= \sum_{t'=t}^T \mathbb{E}_{\pi_\theta} [r(s_{t'}|s_t, a_t)] \\
V^\pi(s_t) &= \mathbb{E}_{a_t \sim \pi_\theta(a_t|s_t)} [Q^\pi(s_t, a_t)] \\
A^\pi(s_t, a_t) &= Q^\pi(s_t, a_t) - V^\pi(s_t)
\end{aligned}$$

1. Take action  $a \sim \pi_\theta(a|s)$ , get  $(s, a, s', r)$ ;
2. Update  $\hat{V}_\phi^\pi$  using  $r + \gamma \hat{V}(s')$
3. Evaluate  $A(s, a) = r(s, a) + \gamma \hat{V}(s') - \hat{V}(s)$
4. PG:  $\nabla_\theta \approx \nabla_\theta \log \pi(a|s) \hat{A}^\pi(s, a)$
5.  $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

What value should we use to multiply with  $\nabla \log \pi(a_t|s_t)$ ?

1. low-variance, biased:  $r(s, a) + \gamma V(s_{t+1}) - V(s_t)$ ;
2. unbiased, high-variance:  $\sum_{t'}^T \gamma^{t'-t} r(s_{t'}, a_{t'}) - b$
3. Eligibility traces and n-step returns: a good balance

$$\hat{A}(s_t, a_t) = \sum_{t'=t}^{t+n} \gamma^{t'-t} r(s_{t'}, a_{t'}) + \gamma^n \hat{V}(s_{t+n}) - \hat{V}(s_t) \quad (9)$$

4. GAE:

$$\hat{A}_{GAE}^\pi(s_t, a_t) = \sum_{n=1}^{\infty} w_n \hat{A}_n^\pi(s_t, a_t) \quad (10)$$

with  $w_n \propto \lambda^{n-1}$

$$\hat{A}_{GAE}^\pi(s_t, a_t) = \sum_{n=1}^{\infty} (\gamma \lambda)^{t'-t} \delta_{t'} \quad (11)$$

## 2.1 System Design

1. Batch-size = num of games (=16 running envs in our scenarios);
2. Policy (actor-critic):
  - (a) Shared CNNBase module: output a  $(?, n)$  dim feature;
  - (b) (Optional)  $x, rnn-hxs = CNNBase.forward-gru(x, rnn-hxs, masks)$
  - (c) Actor head: feature  $\rightarrow (?, n_a)$
  - (d) Critic head: feature  $\rightarrow (?, 1)$
  - (e) Actor head wraps a distribution; for sampling action  $a_t$  and log-prob  $\log \pi_\theta(a_t|s_t)$
3. Rollouts (data pool):
  - (a) Observation:  $(T+1, ?, 4, 84, 84)$
  - (b) RNN hidden state:  $(T+1, ?, nh)$
  - (c) Rewards:  $(T, ?, 1)$
  - (d) Value-preds:  $(T+1, ?, 1)$
  - (e) Returns:  $(T+1, ?, 1)$
  - (f) Action-log-prob:  $(T, ?, 1)$
  - (g) Actions:  $(T, ?, 1)$
  - (h) Masks:  $(T+1, ?, 1)$ ; 1.0 for non-terminal, 0. for terminal;
4. Agent (A2C-ACKTR, gail, kfac, ppo)
  - (a) actor-critic: policy;
  - (b) Coefficients of different terms;
  - (c) Optimizer (KFAC, RMSprop);
  - (d) update();
5. Running Logic:

```
for n steps:
    step 1: for t = 1..5 (no-gradient mode)
        value, act, act-log-prob, rnn-h = actor_critic.act(obs[step], rnn-h, mask[step])
        obs, reward, done, infos = envs.step(action)
        update mask
        rollouts.update: obs, rnn-h, action, action-log-prob,
            value, reward, masks, bad-masks
    step 2: (no-gradient mode)
        next_value = actor_critic.get_value(obs[-1]) # get value for T+1
    step 3: (n-step true target value)
        rollouts.compute_returns(next_value) # (T, ?, 1)
        can also be GAE;
    step 4: update parameter
        value_loss, action_loss, dist_entropy = agent.update()
        1. value, action-log-prob, entropy, rnn-hs = actor_critic.evaluate_actions()
        value, action-feat, rnn-hs = CNNBase(inputs, rnn, masks) # (T * ?, 4, 84, 84)
        dist (action distribution)
        action-log-probs
        dist-entropy
        2. Advantages: A(st, at) = rollouts.returns - values
        3.1 action-loss: -(adv * action-log-prob)
        3.2 value-loss: adv ^ 2
        3.3 entropy: -coeff * dist-entropy
    step 5: rollouts.after_update()
        reset T+1 step obs, rnn-hs, masks, bad-masks to [0]
```

## 3 Value Function

Q-Learning:

1. Collect  $(s, a, s', r)$  with some policy ( **$\epsilon$ -greedy exploration**), could do with **replay buffer**
2.  $y_i \leftarrow r(s_i, a_i) + \gamma \max_a Q(s', a)$
3.  $\phi \leftarrow \arg \max_{\phi'} \sum_i ||Q_{\phi'}(s, a) - y||^2$

Boltzmann exploration:

$$\pi(a_t|s_t) \sim \exp(Q_{\phi}(s_t, a_t))$$

Advanced tricks:

1. Double Q-learning: another network to select  $a'$ :

$$y = r + \gamma Q_{\phi'}(s', \arg \max_{a'} Q_{\phi}(s', a')) \quad (12)$$

2. Q-learning with n-steps: less-biased, typically faster learning especially early on;
3. Continuous actions: DDPG, another network for action  $a$ ;
4. Prioritized experience replay;
5. Clip gradient or Huber loss;